

# CUDA OPTIMIZATION TIPS, TRICKS AND TECHNIQUES

Stephen Jones, GTC 2017




# optimization

*noun* | op·ti·mi·za·tion | \,äp-tə-mə-'zā-shən\


Popularity: Bottom 50% of words

The art of doing **more** with **less**

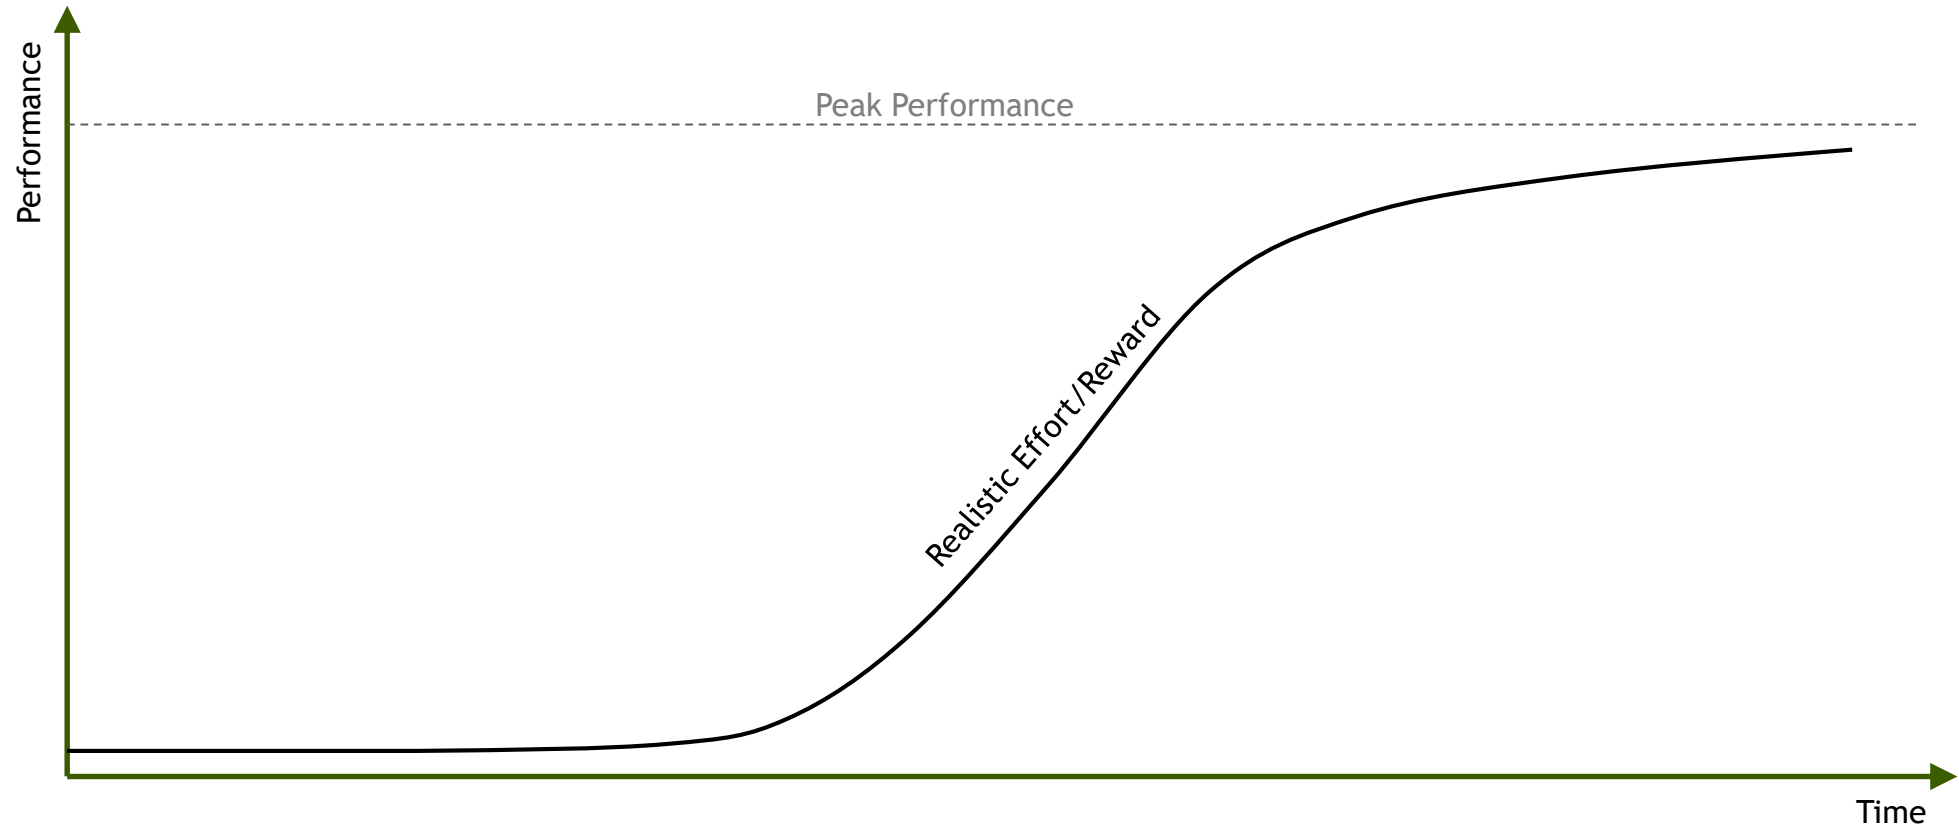
# RULE #1: DON'T TRY TOO HARD




# RULE #1: DON'T TRY TOO HARD




# RULE #1: DON'T TRY TOO HARD



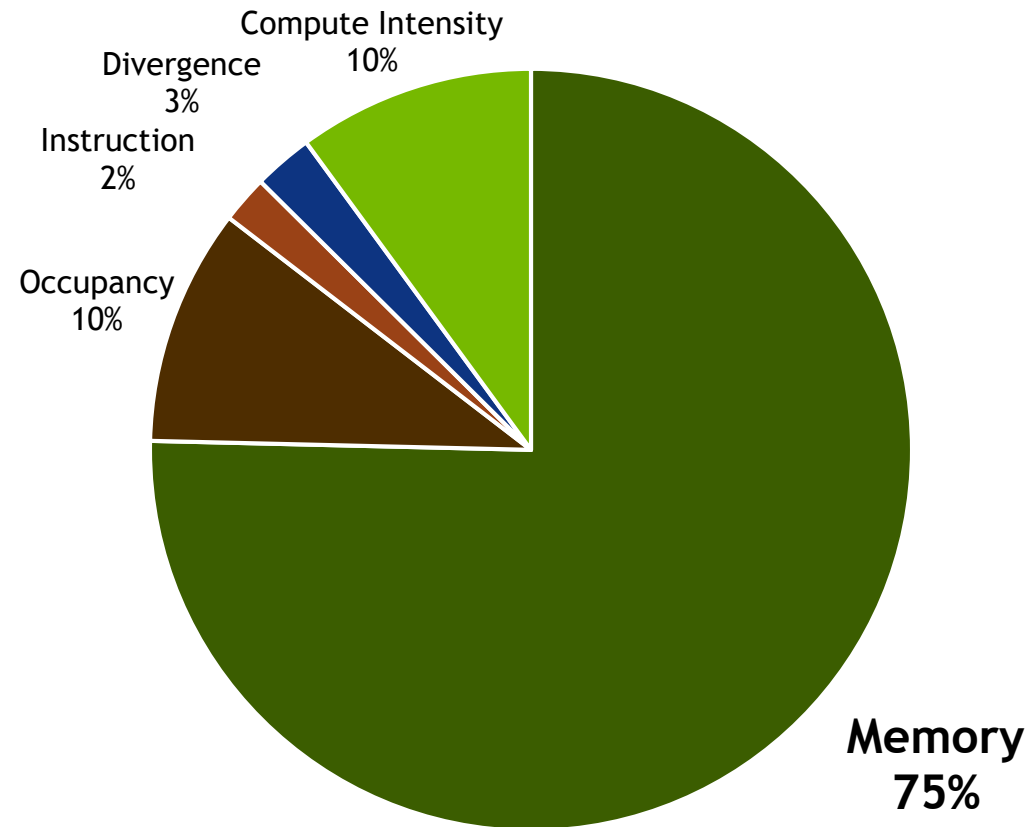
# RULE #1: DON'T TRY TOO HARD



# RULE #1: DON'T TRY TOO HARD




# PERFORMANCE CONSTRAINTS



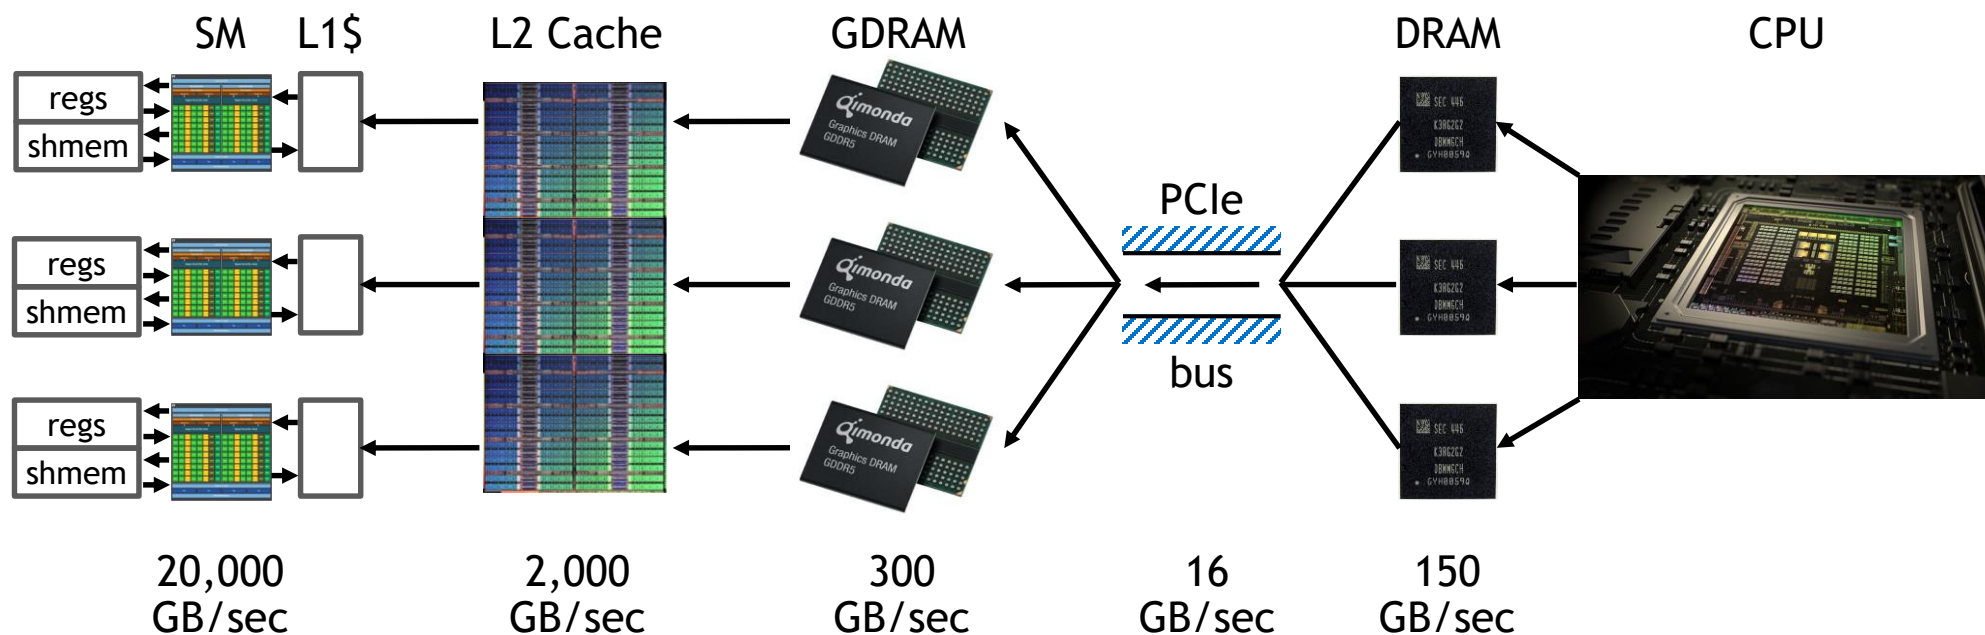


# PERFORMANCE CONSTRAINTS

Chart Title



# MEMORY ORDERS OF MAGNITUDE




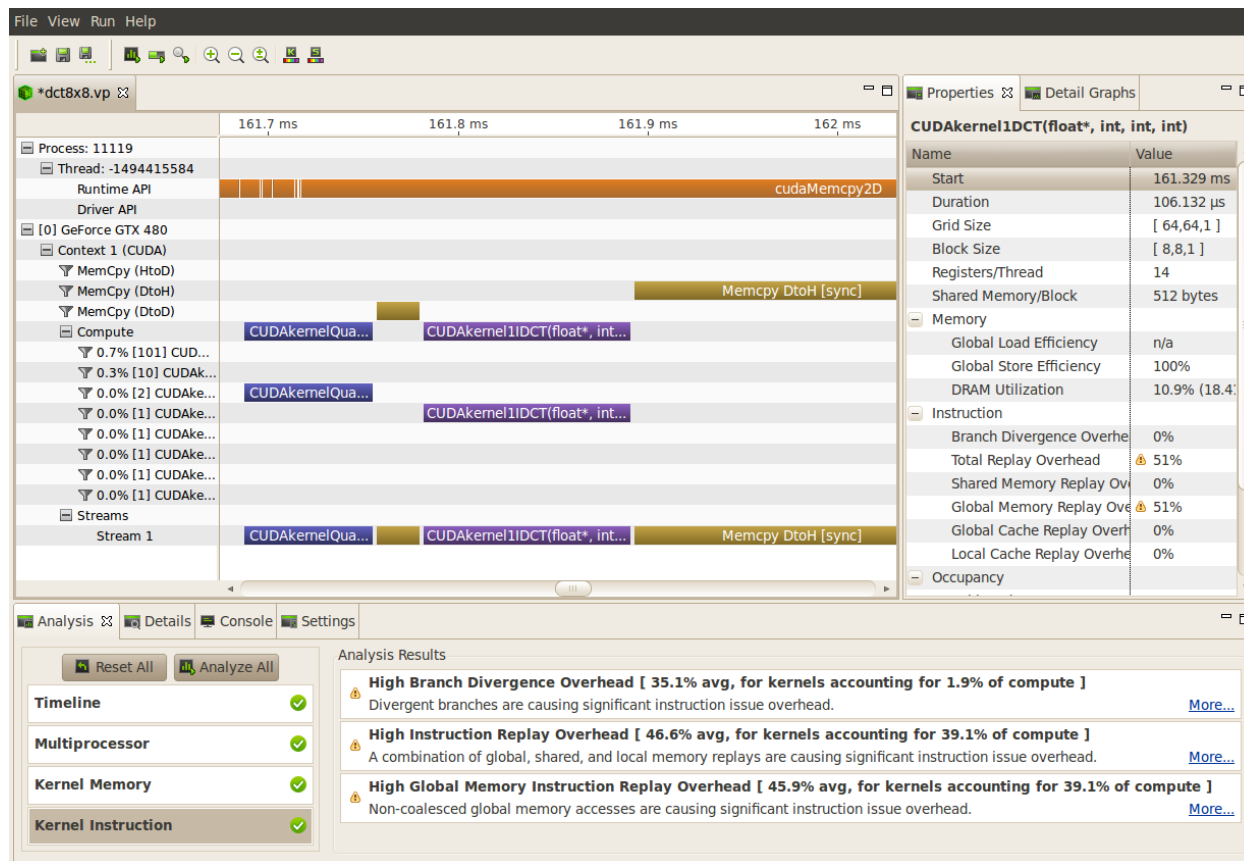
# TALK BREAKDOWN

In no particular order

1. Why Didn't I Think Of That?
2. CPU Memory to GPU Memory (the PCIe Bus)
3. GPU Memory to the SM
4. Registers & Shared Memory
5. Occupancy, Divergence & Latency
6. Weird Things You Never Thought Of (and probably shouldn't try)

**WHERE TO BEGIN?**

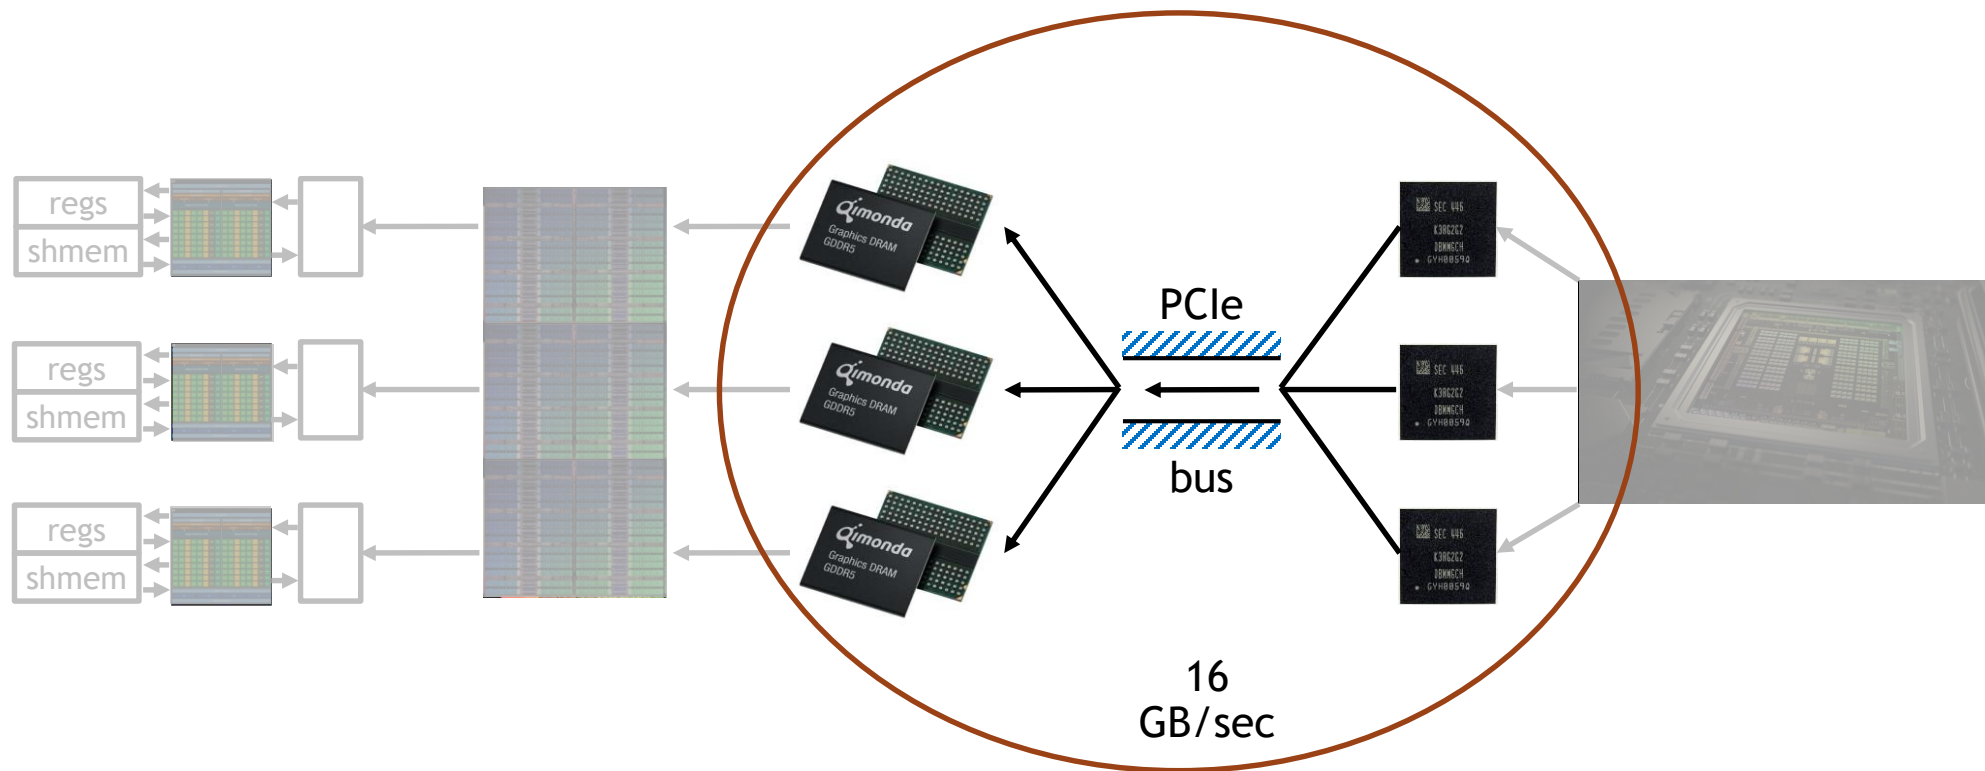
# THE OBVIOUS



Start with the Visual Profiler

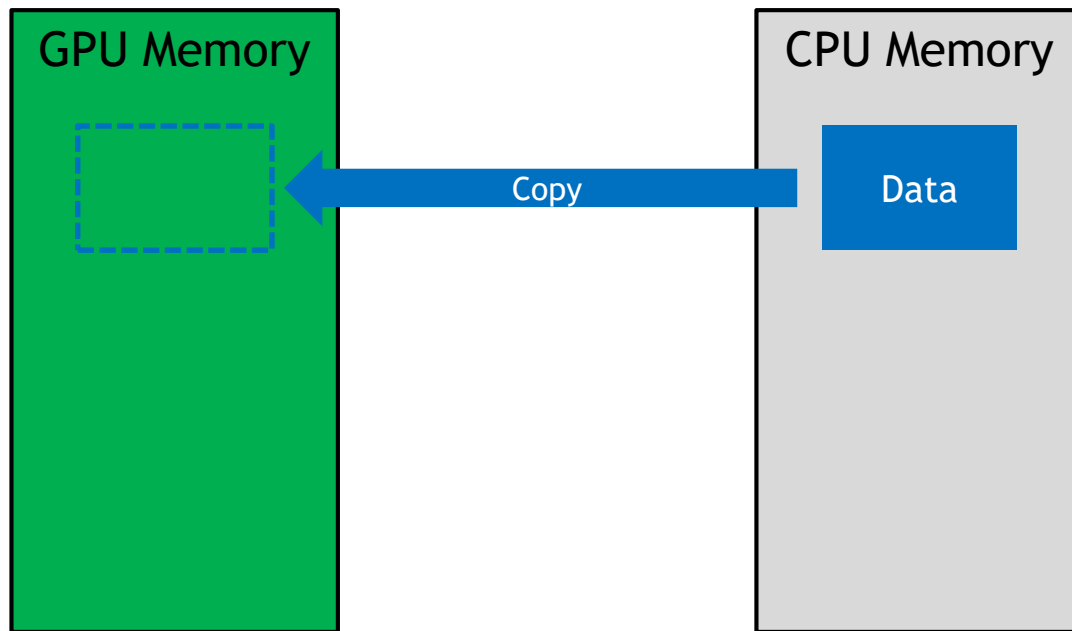
**CPU <> GPU DATA MOVEMENT**

# PCI ISSUES



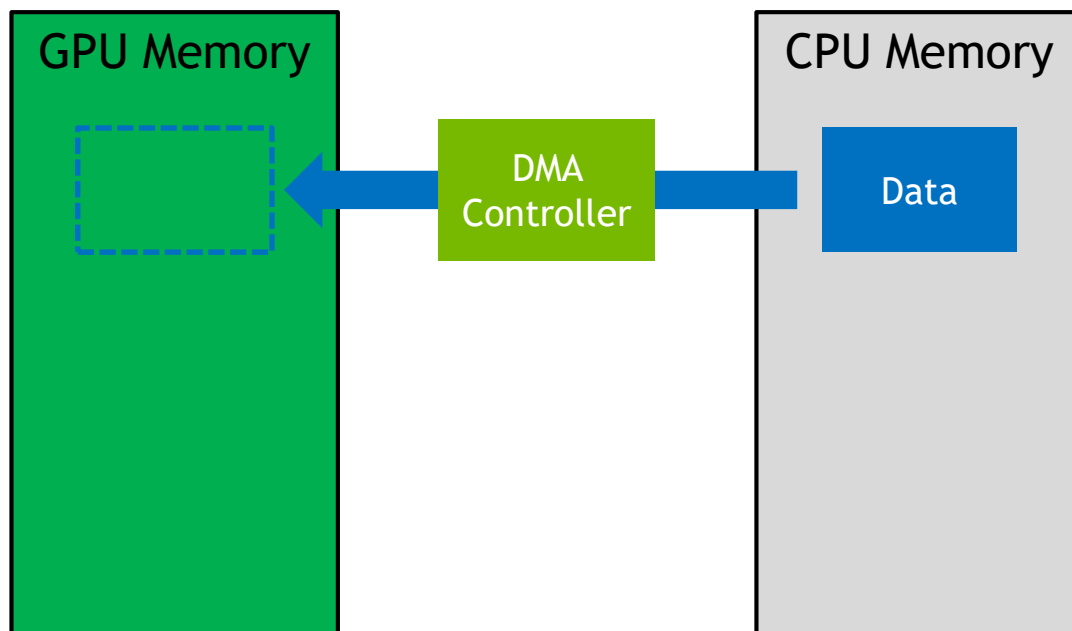
Moving data over the PCIe bus

# PIN YOUR CPU MEMORY

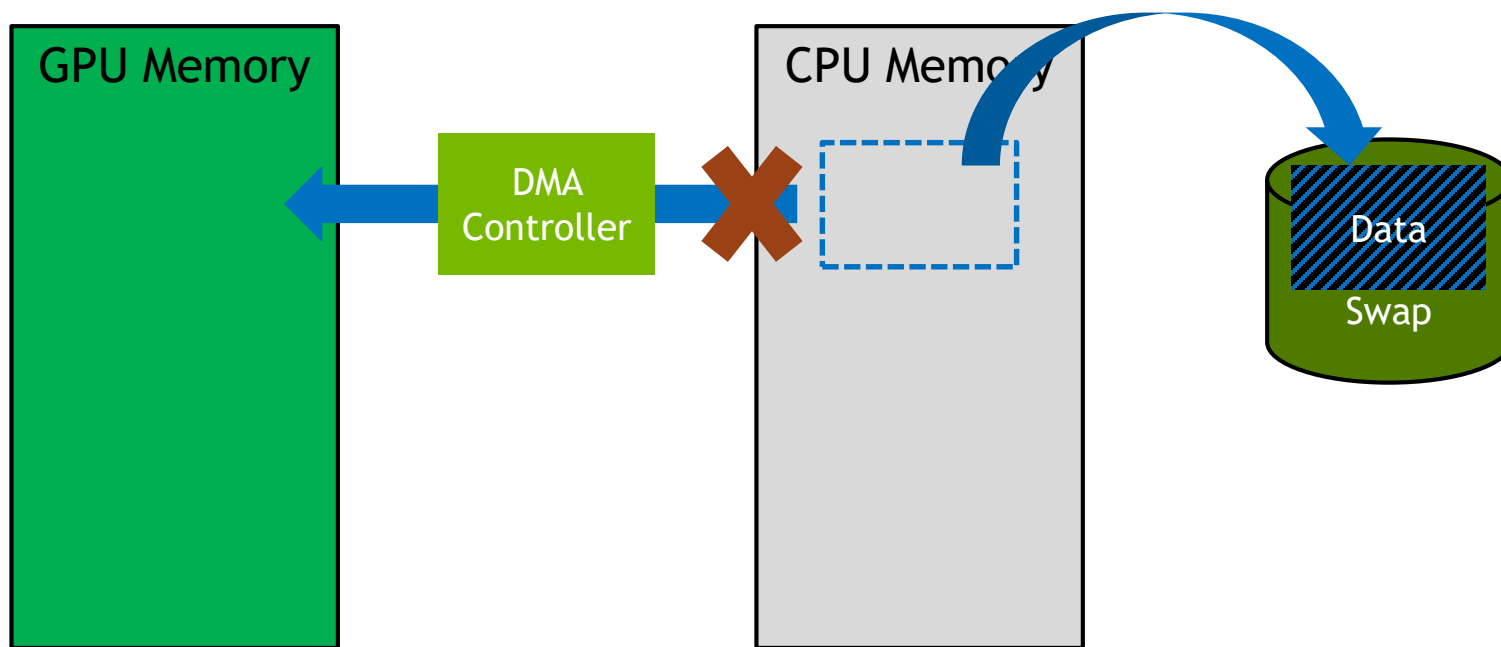




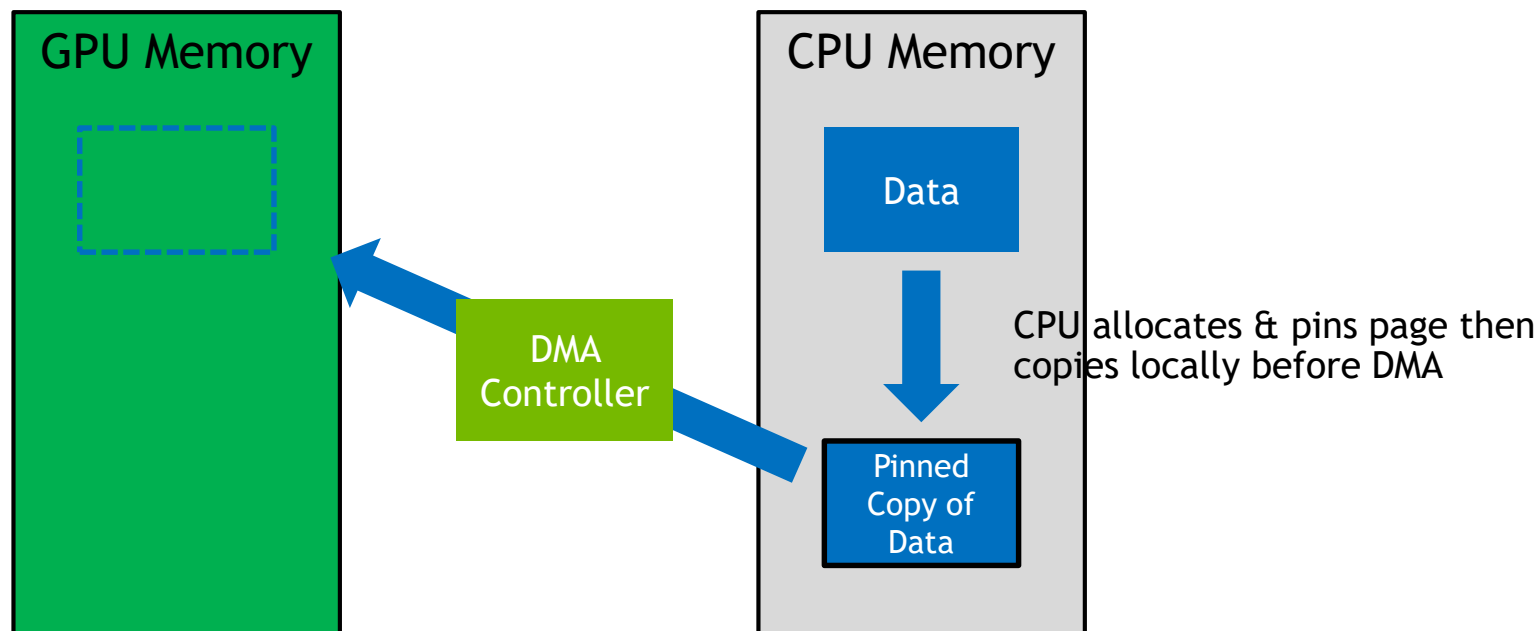
# PIN YOUR CPU MEMORY



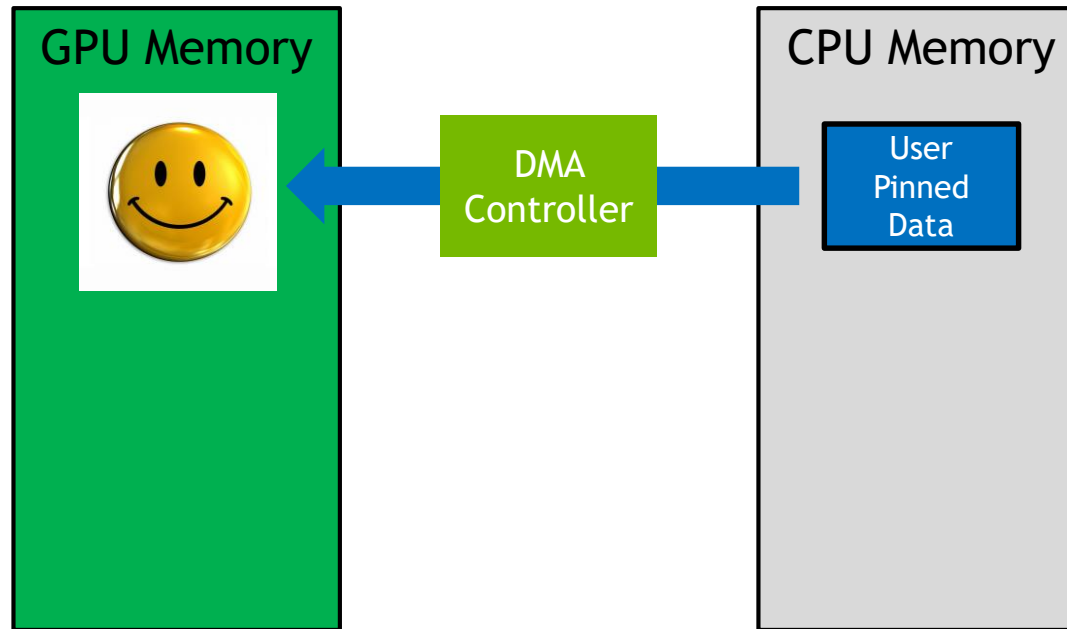
# PIN YOUR CPU MEMORY



# PIN YOUR CPU MEMORY




# PIN YOUR CPU MEMORY




```
cudaHostAlloc( &data, size, cudaHostAllocMapped );  
cudaHostRegister( &data, size, cudaHostRegisterDefault );
```

# PIN YOUR CPU MEMORY



# REMEMBER: PCIe GOES BOTH WAYS

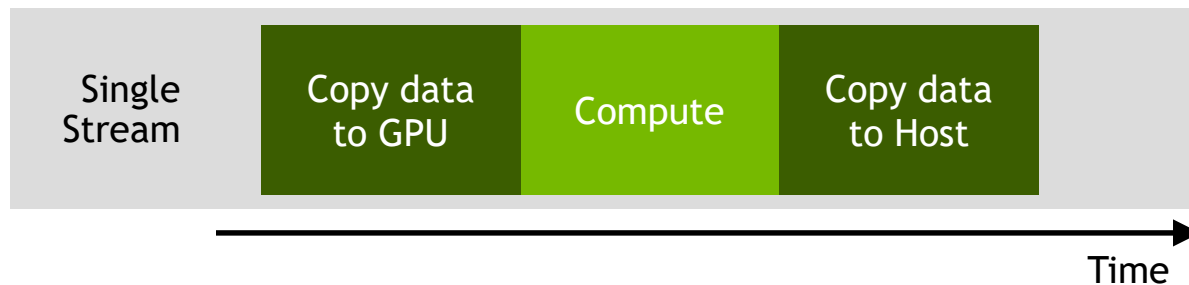


# STREAMS & CONCURRENCY

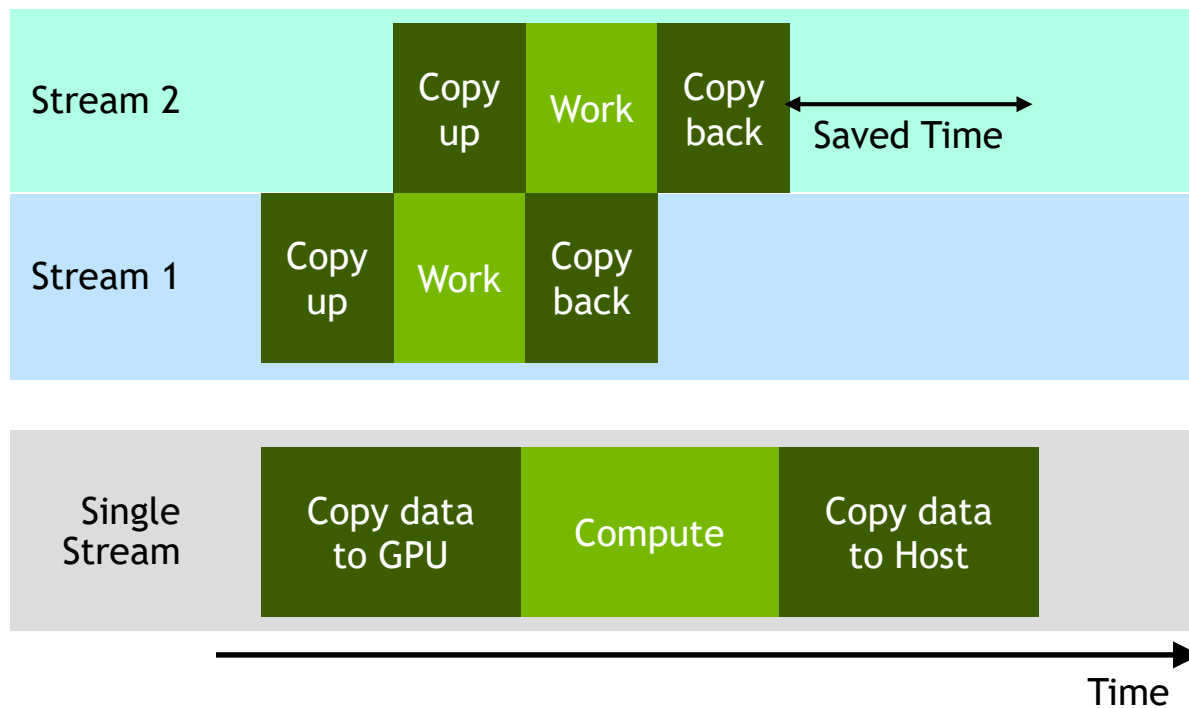
Hiding the cost of data transfer

Operations in a single stream are ordered

But hardware can copy and compute at the same time



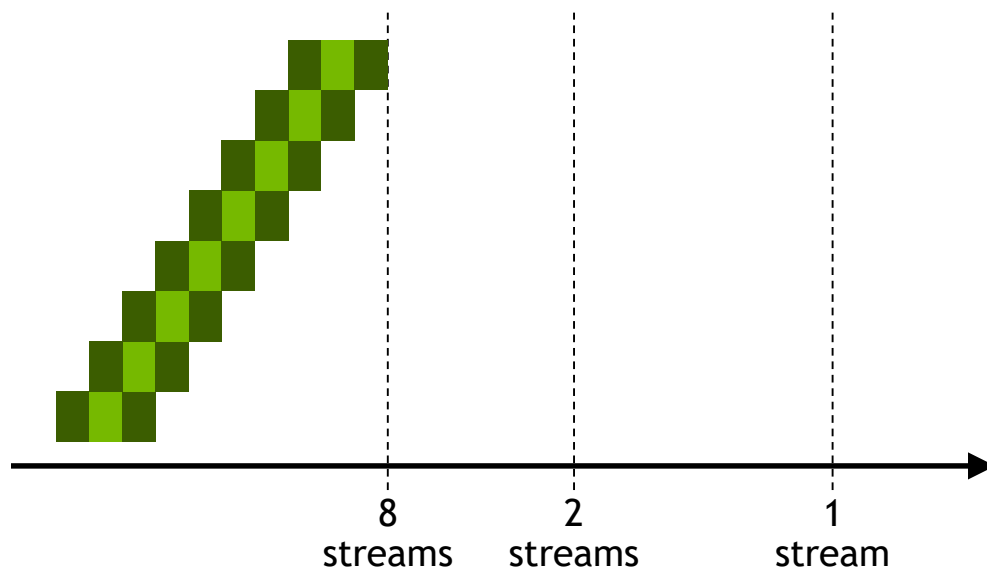
# STREAMS & CONCURRENCY





# STREAMS & CONCURRENCY

Can keep on breaking work into smaller chunks and saving time

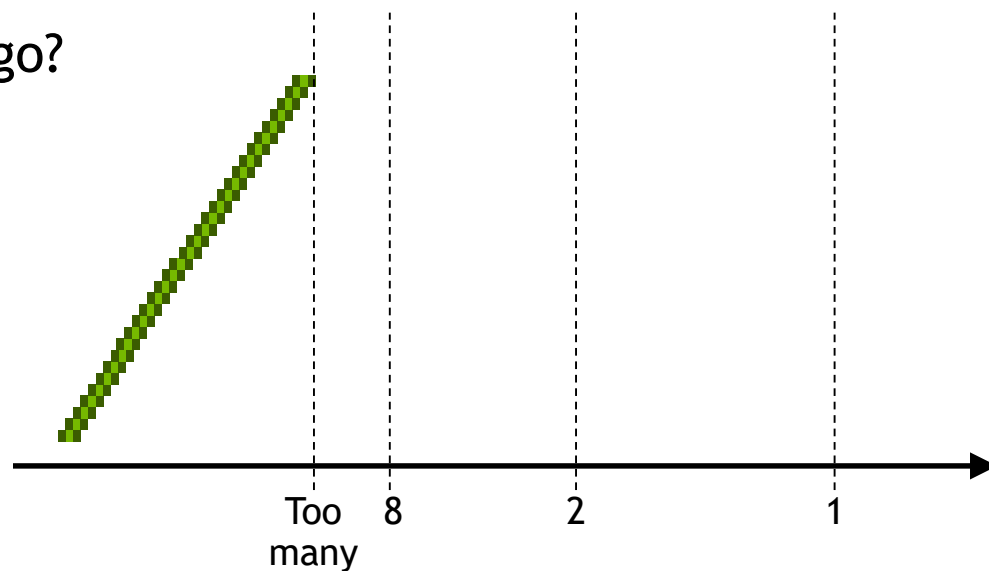


# SMALL PCIe TRANSFERS


PCIe is designed for large data transfers

But fine-grained copy/compute overlap prefers small transfers

So how small can we go?

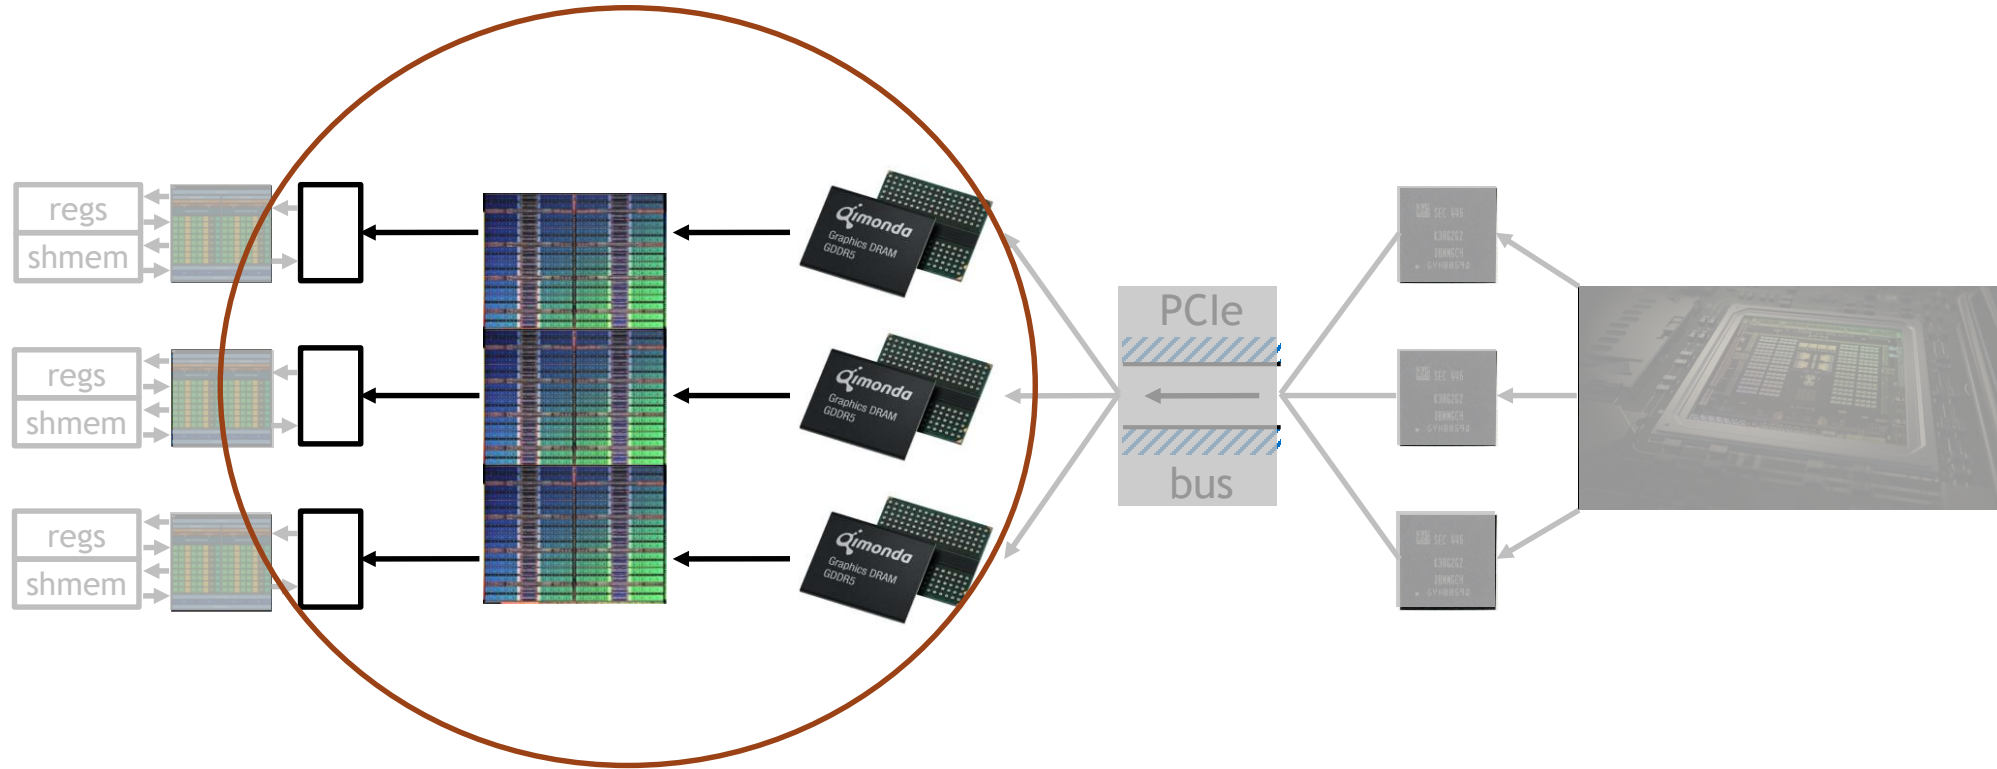


# APPARENTLY NOT THAT SMALL



# FROM GPU MEMORY TO GPU THREADS

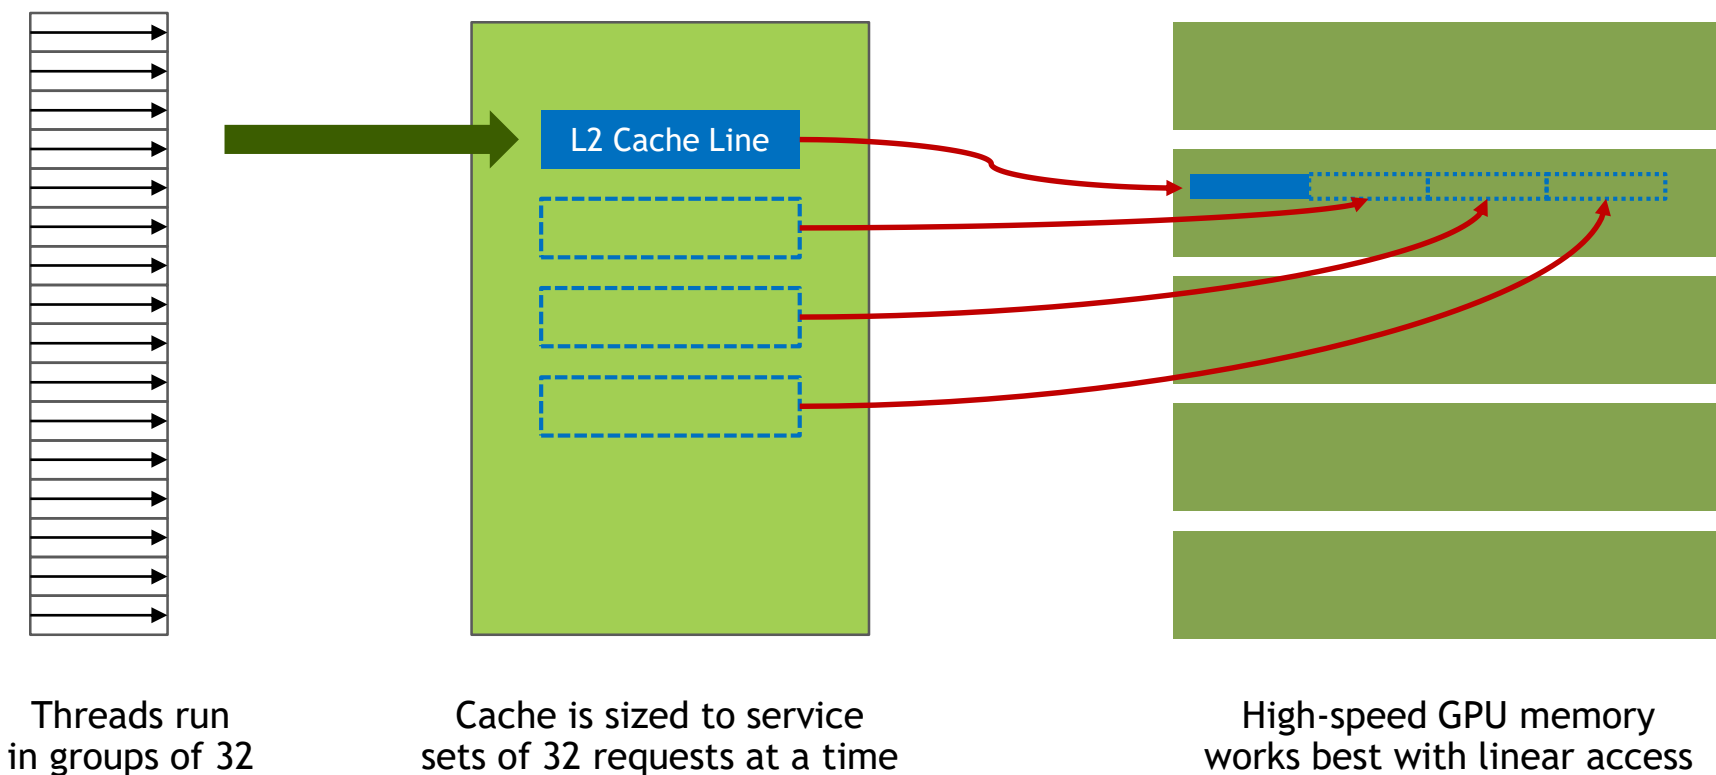
# FEEDING THE MACHINE



From GPU Memory to the SMs


# USE THE PARALLEL ARCHITECTURE

Hardware is optimized to use all SIMT threads at once




# VECTORIZE MEMORY LOADS

Multi-Word as well as Multi-Thread




# VECTORIZIZE MEMORY LOADS



Fill multiple cache lines in a single fetch



# VECTORIZIZE MEMORY LOADS



Fill multiple cache lines in a single fetch

# VECTORIZIZE MEMORY LOADS



# DO MULTIPLE LOADS PER THREAD

Multi-Thread, Multi-Word AND Multi-Iteration

```
__global__ void copy(int2 *input,
                    int2 *output,
                    int max) {

    int id = threadIdx.x +
            blockDim.x * blockIdx.x;

    if( id < max ) {
        output[id] = input[id];
    }
}
```

One copy per thread  
Maximum overhead

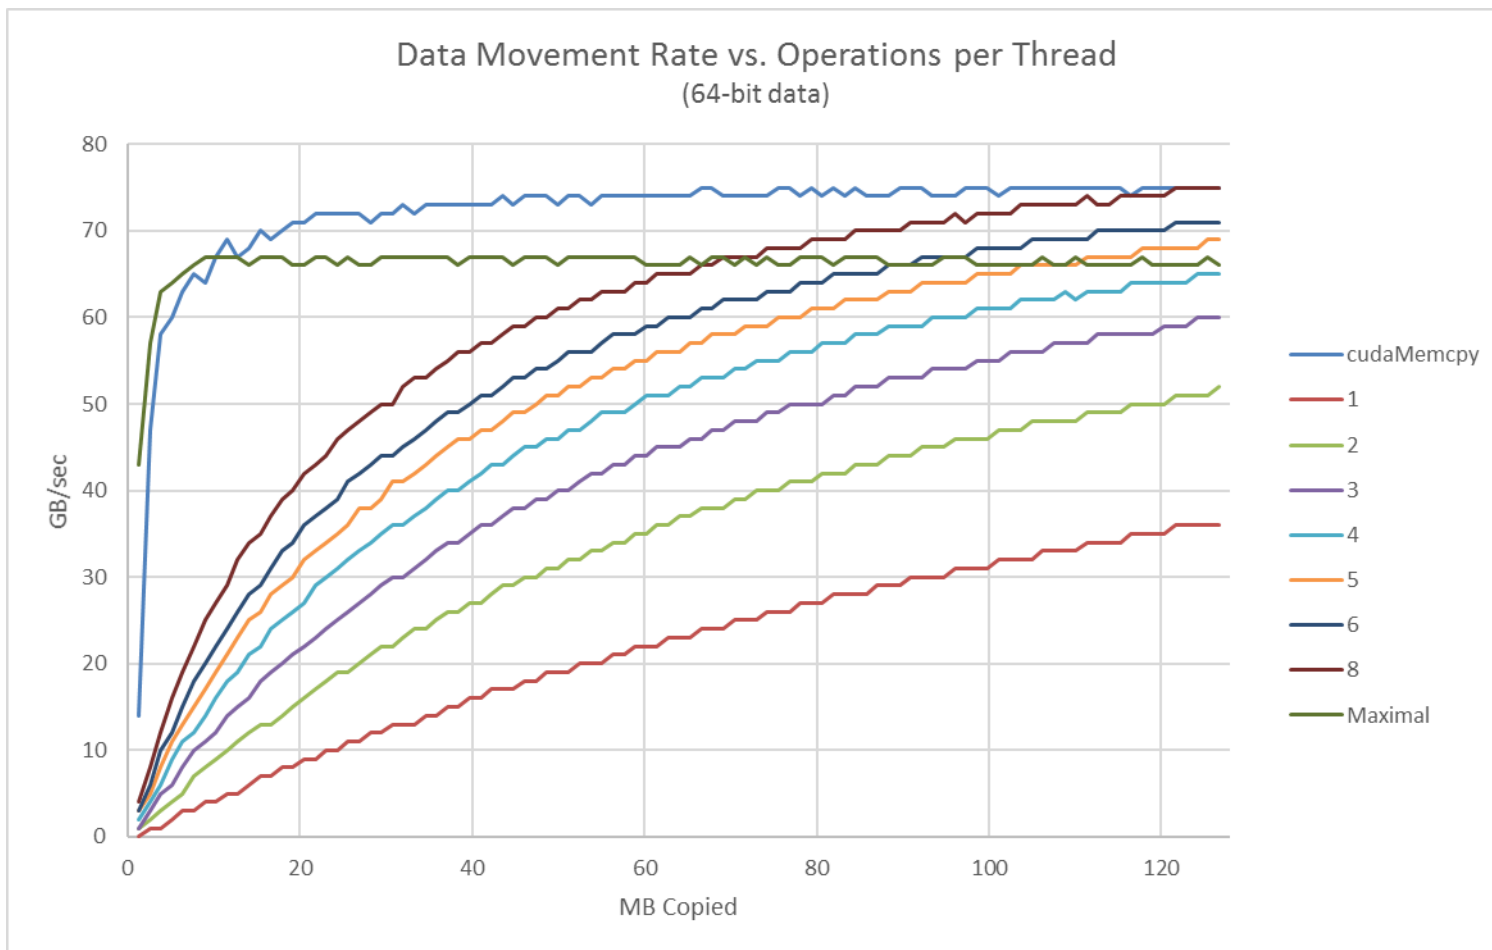
```
__global__ void copy(int2 *input,
                    int2 *output,
                    int max,
                    int loadsPerThread) {

    int id = threadIdx.x +
            blockDim.x * blockIdx.x;

    for(int n=0; n<loadsPerThread; n++) {
        if( id >= max ) {
            break;
        }
        output[id] = input[id];
        id += blockDim.x * gridDim.x;
    }
}
```

Multiple copies per thread  
Amortize overhead


# “MAXIMAL” LAUNCHES ARE BEST




# COALESCED MEMORY ACCESS

It's not just good enough to use all SIMT threads


Coalesced: Sequential memory accesses are adjacent




Uncoalesced: Sequential memory accesses are unassociated




# SIMT PENALTIES WHEN NOT COALESCED





Single 32-wide operation




32 one-wide operations



# SCATTER & GATHER




# AVOID SCATTER/GATHER IF YOU CAN



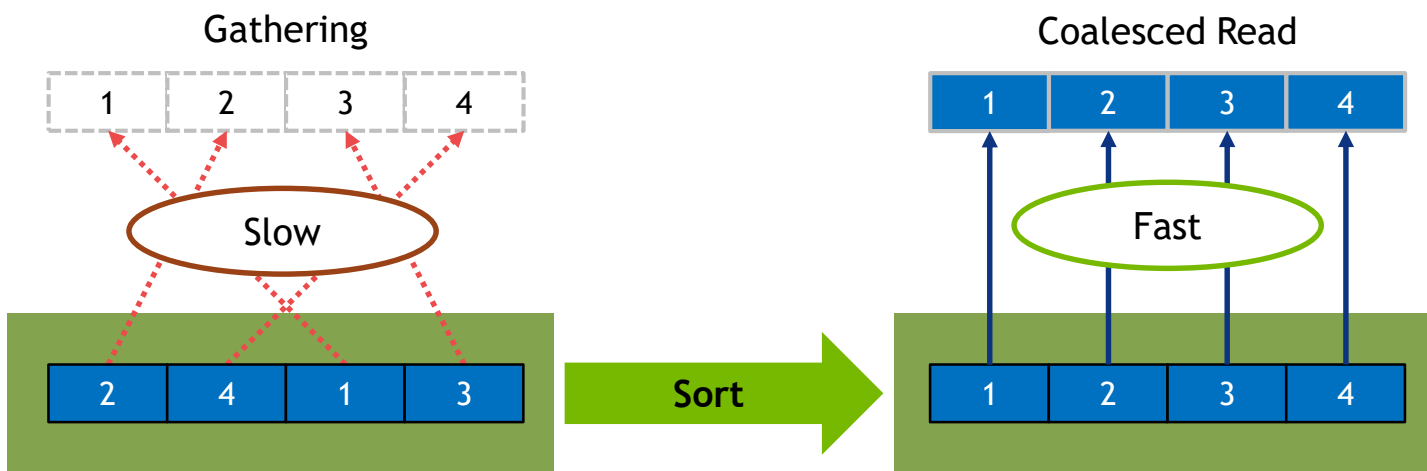


# AVOID SCATTER/GATHER IF YOU CAN




# SORTING MIGHT BE AN OPTION

If reading non-sequential data is expensive, is it worth sorting it to make it sequential?




# SORTING MIGHT BE AN OPTION

Even if you're only going to read it twice, then yes!



# PRE-SORTING TURNS OUT TO BE GOOD



# DATA LAYOUT: “AOS vs. SOA”

Sometimes you can't just sort your data

## Array-of-Structures

```
#define NPTS 1024 * 1024

struct Coefficients_AOS {
    double u[3];
    double x[3][3];
    double p;
    double rho;
    double eta;
};

Coefficients_AOS gridData[NPTS];
```

Single-thread code prefers arrays of structures, for cache efficiency

## Structure-of-Arrays

```
#define NPTS 1024 * 1024

struct Coefficients_SOA {
    double u[3][NPTS];
    double x[3][3][NPTS];
    double p[NPTS];
    double rho[NPTS];
    double eta[NPTS];
};

Coefficients_SOA gridData;
```

SIMT code prefers structures of arrays, for execution & memory efficiency


# DATA LAYOUT: “AOS vs. SOA”

```
#define NPTS 1024 * 1024

struct Coefficients_AOS {
    double u[3];
    double x[3][3];
    double p;
    double rho;
    double eta;
};

Coefficients_AOS gridData[NPTS];
```

Structure Definition




Conceptual Layout

# SOA: STRIDED ARRAY ACCESS

GPU reads data one element at a time, but in parallel by 32 threads in a warp

u0	u1	u2
x00	x01	x02
x10	x11	x12
x20	x21	x22
p		
rho		
eta		


Conceptual Layout



```
double u0 = gridData[threadIdx.x].u[0];
```

# AOS: COALESCED BUT COMPLEX

GPU reads data one element at a time, but in parallel by 32 threads in a warp




Conceptual Layout

Array-of-Structures Memory Layout



Structure-of-Arrays Memory Layout




```
double u0 = gridData.u[0][threadIdx.x];
```



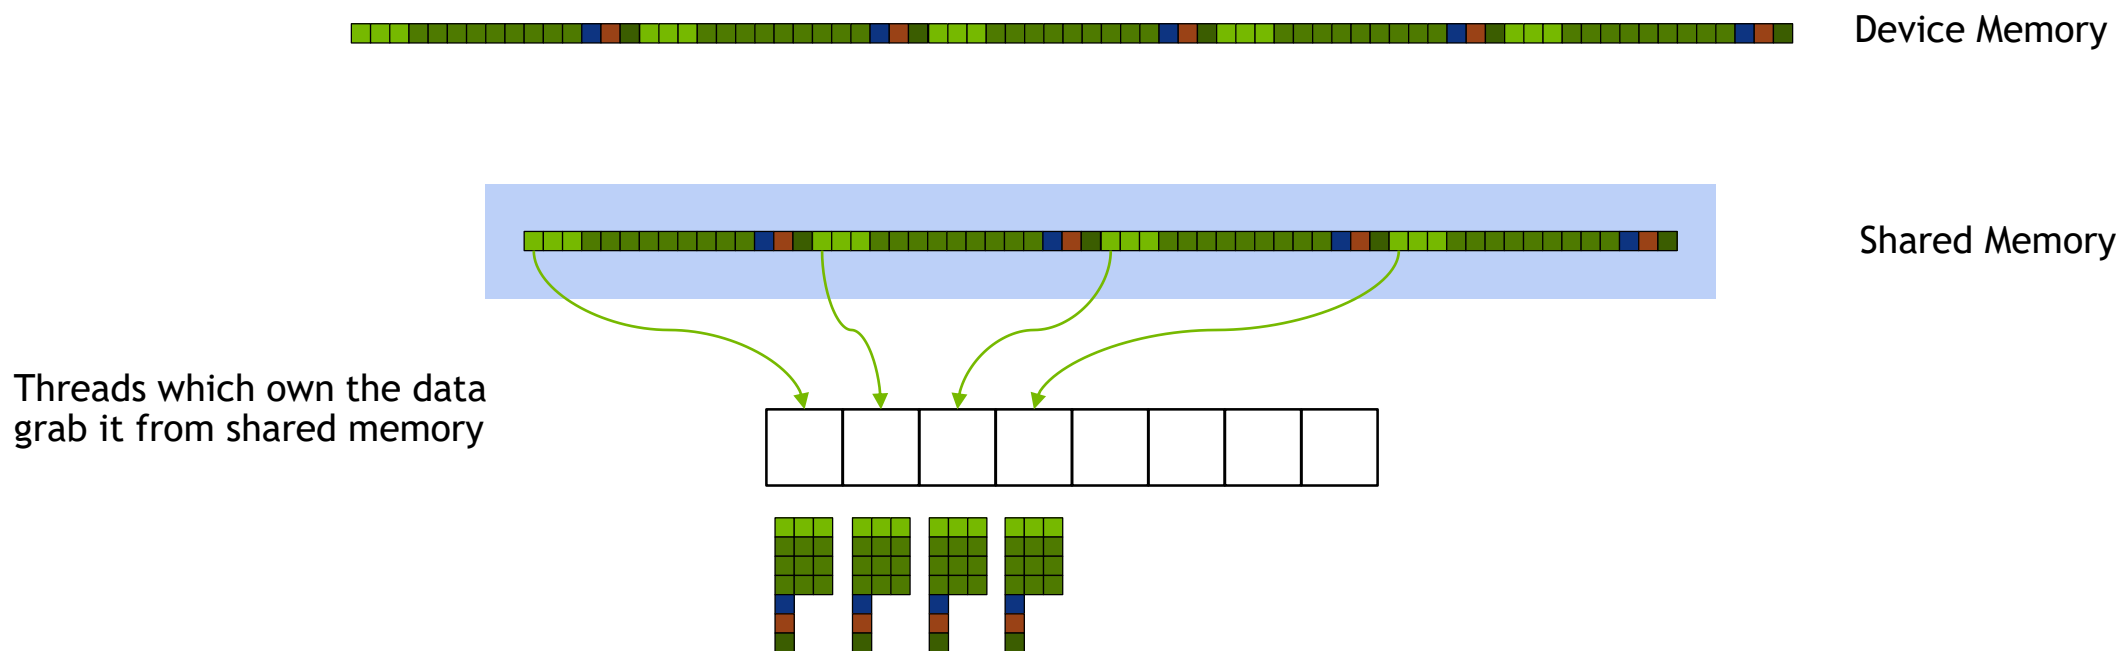
# BLOCK-WIDE LOAD VIA SHARED MEMORY

Read data linearly as bytes. Use shared memory to convert to struct




# BLOCK-WIDE LOAD VIA SHARED MEMORY

Read data linearly as bytes. Use shared memory to convert to struct



# CLEVER AOS/SOA TRICKS



# CLEVER AOS/SOA TRICKS



Helps for any data size



# HANDY LIBRARY TO HELP YOU

Trove - A utility library for fast AOS/SOA access and transposition

<https://github.com/bryancatanzaro/trove>




**(AB)USING THE CACHE**


# MAKING THE MOST OF L2-CACHE

L2 cache is **fast** but small:

Architecture	L2 Cache Size	Total Threads	Cache Bytes per Thread
Kepler	1536 KB	30,720	51
Maxwell	3072 KB	49,152	64
Pascal	4096 KB	114,688	36

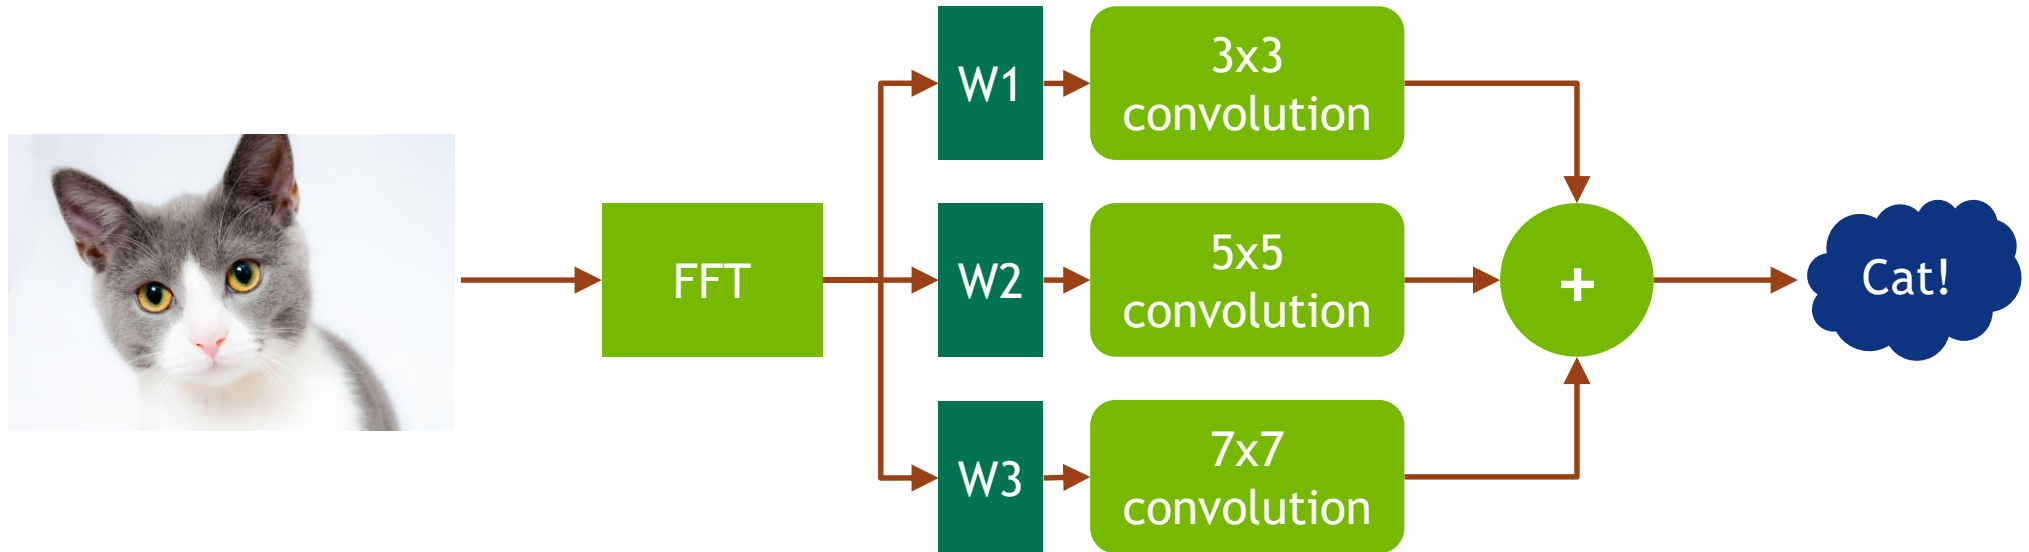


# TRAINING DEEP NEURAL NETWORKS

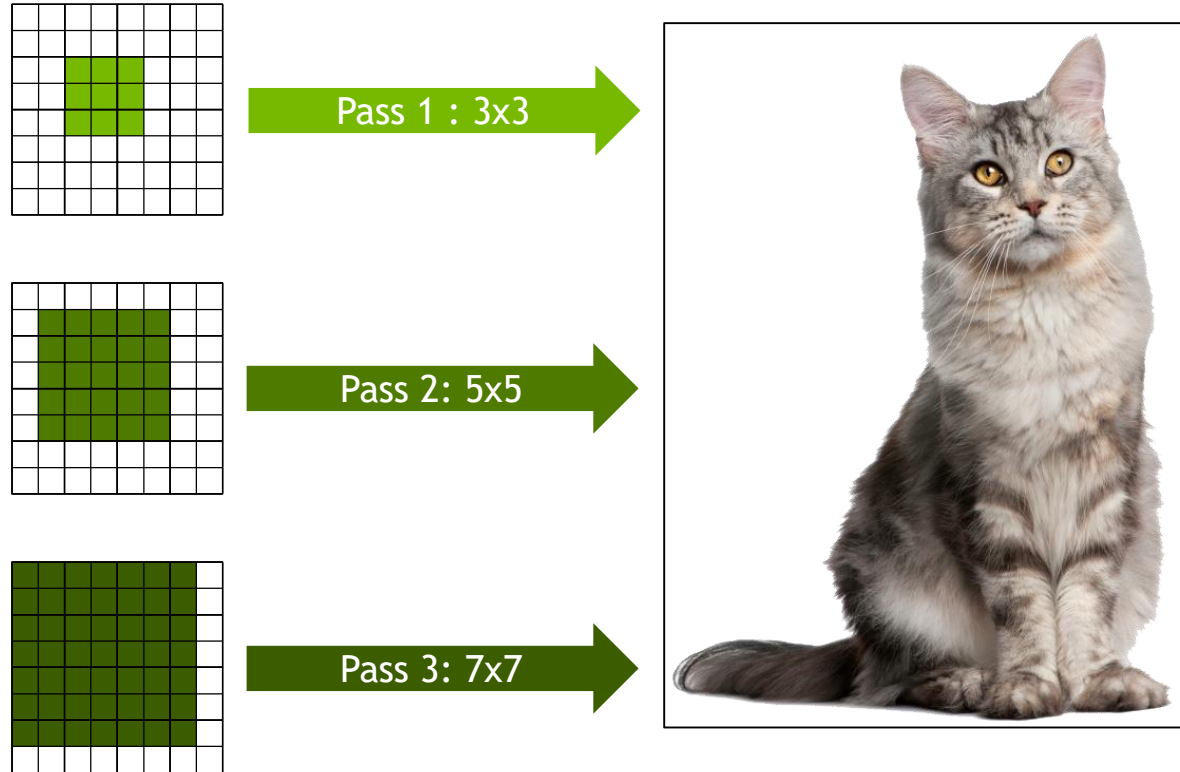








# LOTS OF PASSES OVER DATA



# MULTI-RESOLUTION CONVOLUTIONS




# TILED, MULTI-RESOLUTION CONVOLUTION



Do 3 passes per-tile

Each tile sized to fit in L2 cache

# LAUNCHING FEWER THAN MAXIMUM THREADS




**SHARED MEMORY: DEFINITELY WORTH IT**

# USING SHARED MEMORY WISELY

Shared memory arranged into “banks” for concurrent SIMT access

- 32 threads can read simultaneously so long as into separate banks

Shared memory has 4-byte and 8-byte “bank” sizes



# STENCIL ALGORITHM

Many algorithms have high data re-use: potentially good for shared memory




“Stencil” algorithms accumulate data from neighbours onto a central point

- Stencil has width “ $W$ ” (in the above case,  $W=5$ )


Adjacent threads will share  $(W-1)$  items of data - good potential for data re-use

# STENCILS IN SHARED MEMORY





# SIZE MATTERS



# PERSISTENT KERNELS

## Revisiting the tiled convolutions

Avoid multiple kernel launches by caching in shared memory instead of L2


```
void tiledConvolution() {  
    convolution<3><<< numblocks, blockdim, 0, s >>>(ptr, chunkSize);  
    convolution<5><<< numblocks, blockdim, 0, s >>>(ptr, chunkSize);  
    convolution<7><<< numblocks, blockdim, 0, s >>>(ptr, chunkSize);  
}
```

Separate kernel  
launches with  
L2 re-use

```
__global__ void convolutionShared(int *data, int count, int sharedelems) {  
    extern __shared__ int shdata[];  
    shdata[threadIdx.x] = data[threadIdx.x + blockDim.x*blockIdx.x];  
    __syncthreads();  
  
    convolve<3>(threadIdx.x, shdata, sharedelems);  
    __syncthreads();  
    convolve<5>(threadIdx.x, shdata, sharedelems);  
    __syncthreads();  
    convolve<7>(threadIdx.x, shdata, sharedelems);  
}
```

Single kernel  
launch with  
persistent kernel

# PERSISTENT KERNELS



# OPERATING DIRECTLY FROM CPU MEMORY

Can save memory copies. It's obvious when you think about it ...




Compute only begins when 1<sup>st</sup> copy has finished. Task only ends when 2<sup>nd</sup> copy has finished.



Compute begins after first fetch. Uses lots of threads to cover host-memory access latency. Takes advantage of bi-directional PCI.

# OPERATING DIRECTLY FROM CPU MEMORY



# OCCUPANCY AND REGISTER LIMITATIONS

Register file is bigger than shared memory and L1 cache!

Occupancy can kill you if you use too many registers

Often worth forcing fewer registers to allow more blocks per SM

But watch out for math functions!

```
__launch_bounds__(maxThreadsPerBlock,  
                  minBlocksPerMultiprocessor)  
  
__global__ void compute() {  
    y = acos(pow(log(fdivide(tan(cosh(erfc(x)))), 2)), 3);  
}
```

Function	float	double
log	7	18
cos	16	28
acos	6	18
cosh	7	10
tan	15	28
erfc	14	22
exp	7	10
log10	6	18
normcdf	16	26
cbrt	8	20
sqrt	6	12
rsqrt	5	12
y0	20	30
y1	22	30
fdivide	11	20
pow	11	24
grad. desc.	14	22

# THANK YOU!

